


Try your best, on #1.
#2 is hard.





Ridge Regression

Alvin Wan
alvinwan.com/cs189





Quote of the day



"Give me your tired, your poor, your huddled
masses yearning to breathe free, the
wretched refuse of your teeming shore."

-- Sinho, EE 126 (to me)






Let's talk content.




1 - Proxy for “True” Error

How well our model represents the true model


True error?




True model




True model




True data




Noise



Observations



Proposed model



True error



Some of that error is irreducible. Inherent
noise in observations.

Why is it irreducible??





Math

Math

Math

Math

Math

Math

Math

Math



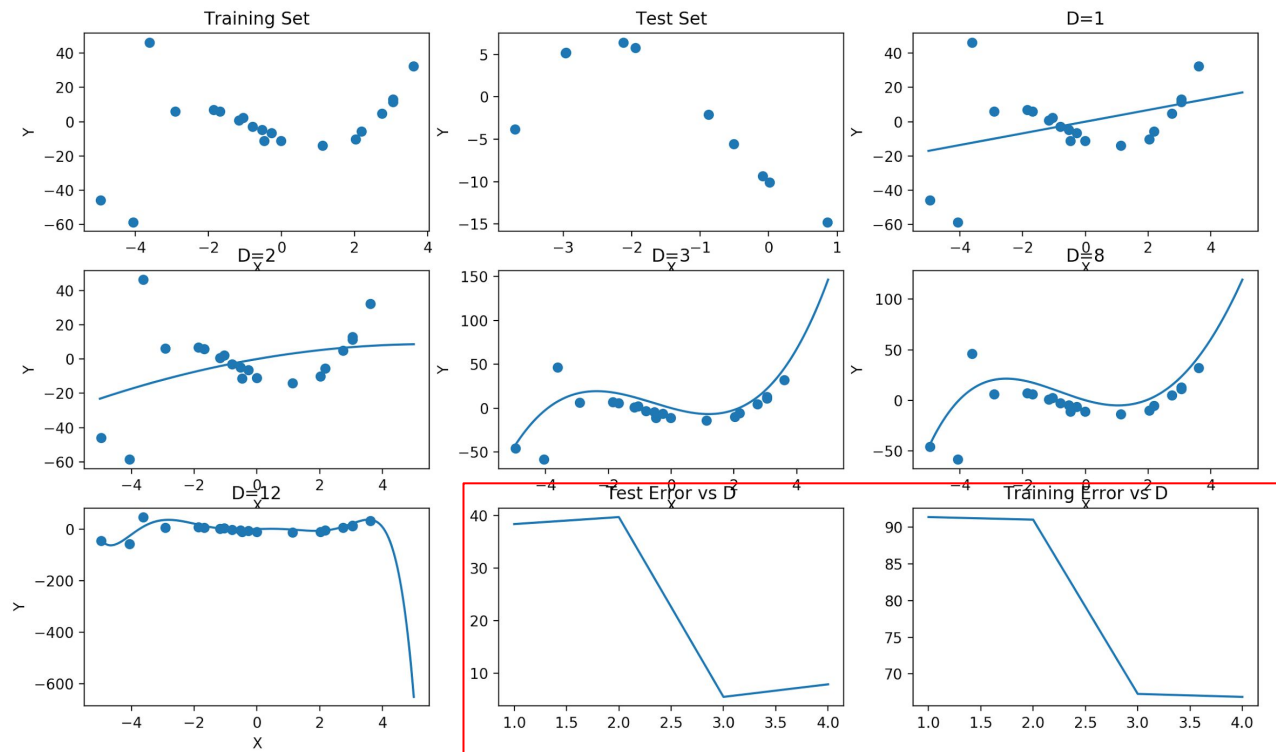
Next week.



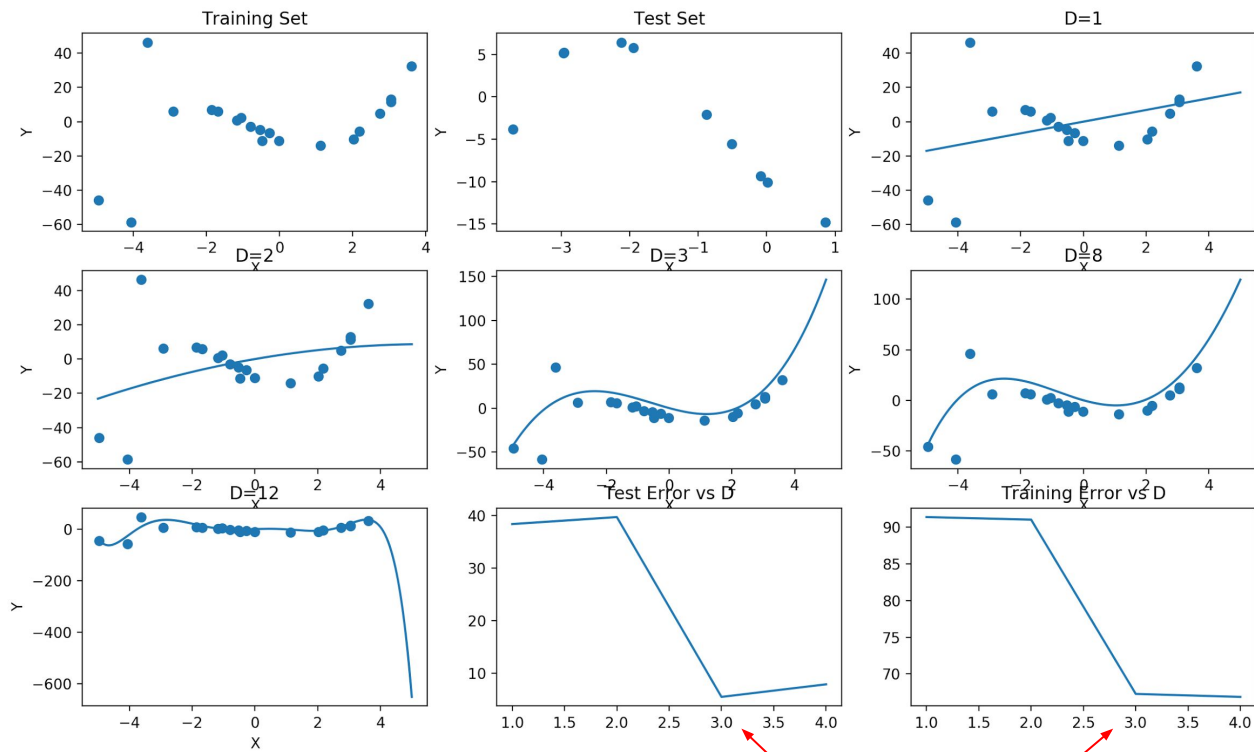


So, our model can suck.



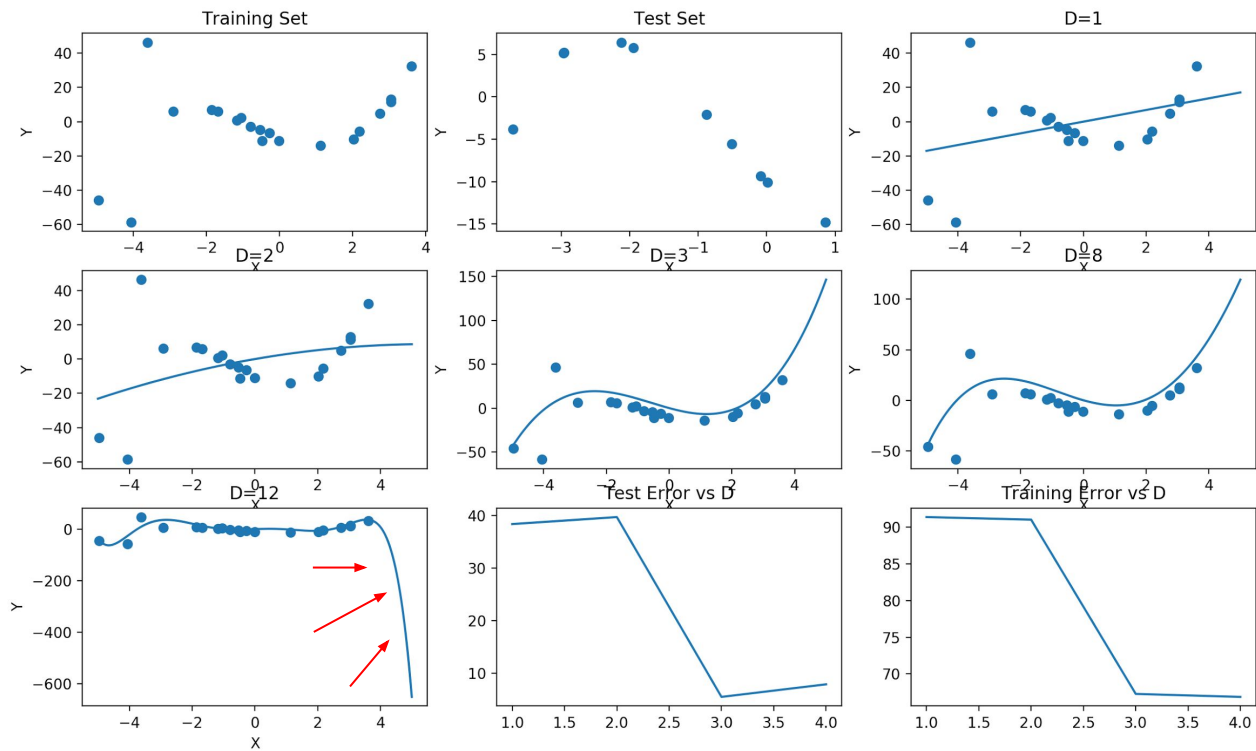


Model order complexity

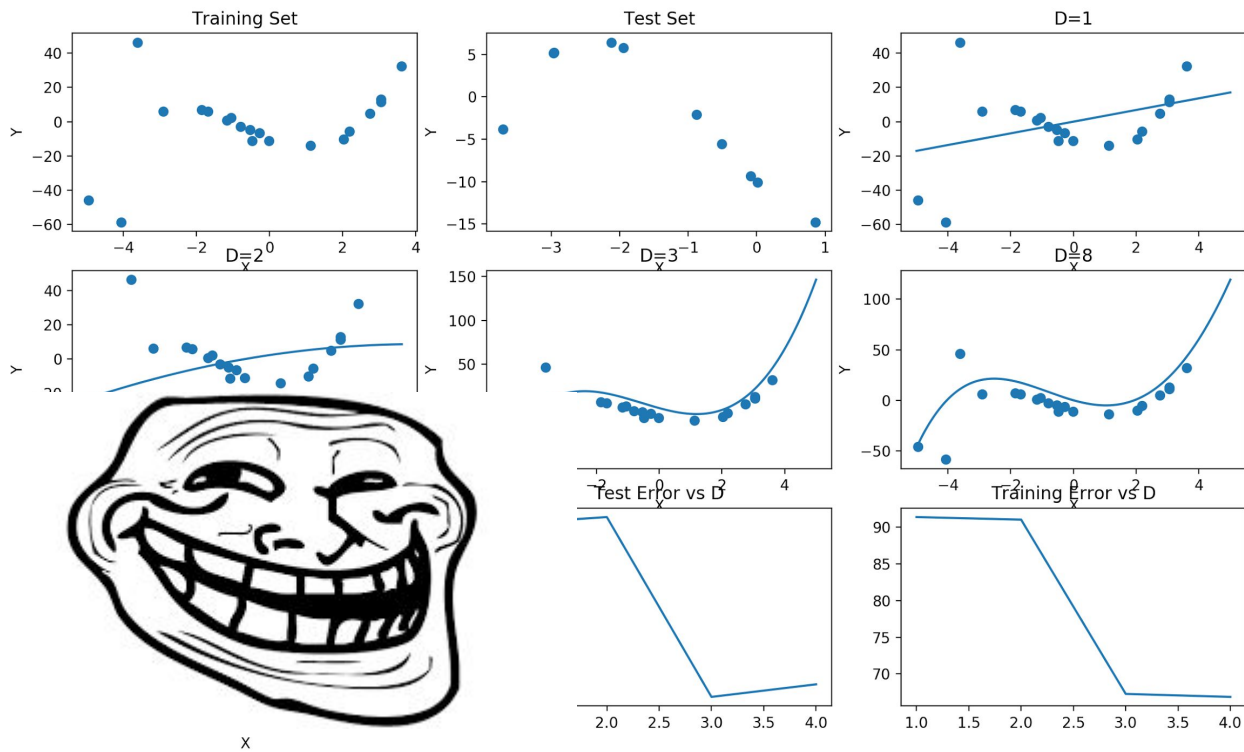


Model order complexity

Sweet spot



Model order complexity



Model order complexity



Overfitting is the enemy...



2 - Tricks of the Trade


Intuition behind regularization



Step 2

Try some matplotlib. Use ``gen_data``







```
data = gen_data(1, 20, 10)
plt.scatter(data['Xtrain'], data['Ytrain'])
```



Step 3

Use ``my_regression``. Play with values.





```
my_regression(data,3,0.1, plot=True)
```



Step 4

Plot as a function of λ .





Step 5

Trick 1: Make your train-val split



```
indices = list(range(n_total))
random.shuffle(indices)
all_x = all_x[indices]
all_y = all_y[indices]
arr_train_x = all_x[: Ntrain]
arr_train_y = all_y[: Ntrain]
arr_val_x = all_x[Ntrain :]
arr_val_y = all_y[Ntrain :]
```




Step 6

Play with different splits, parameters.





Step 7 & 8

Plot as function of model complexity, λ .





Feedback? aaalv.in/survey

